

Assignment4: Deepspeech2

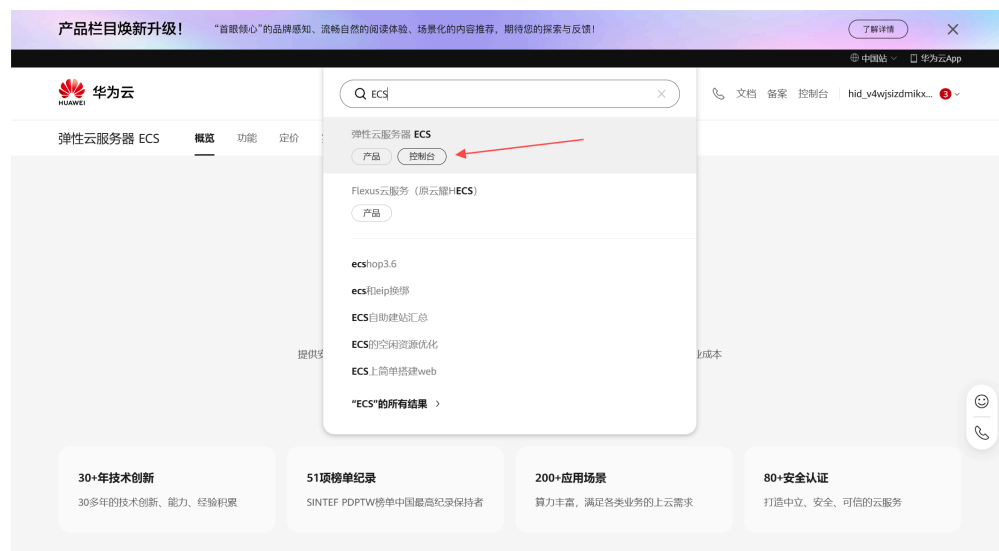
Assignment4: Deepspeech2

1. Development Environment Construction
 - 1.1. Server Purchase
 - 1.2. MobaXterm Connect ECS
2. Code and Data Download
 - 2.1. Get Code Files
 - 2.2. Dataset and its Preprocessing
 - 2.2.1. Download the LibriSpeech Dataset
 - 2.2.2. Install the Python3.9.0
 - 2.2.3. Install the MindSpore and the Required Dependency Package
 - 2.2.4. The Data Preprocessing SeanNaren Scripts was Downloaded
 - 2.2.5. LibriSpeech Data Preprocessing
3. Model Training Results and Evaluation
 - 3.1. Model Training
 - 3.2. Model Training Results and Evaluation
 - 3.2.1. View the Training Log
 - 3.2.2. Training Results and Evaluation
 - 3.3. Model Export
4. Answer the Question
 - 4.1. Question Description
 - 4.2. Problem Thinking

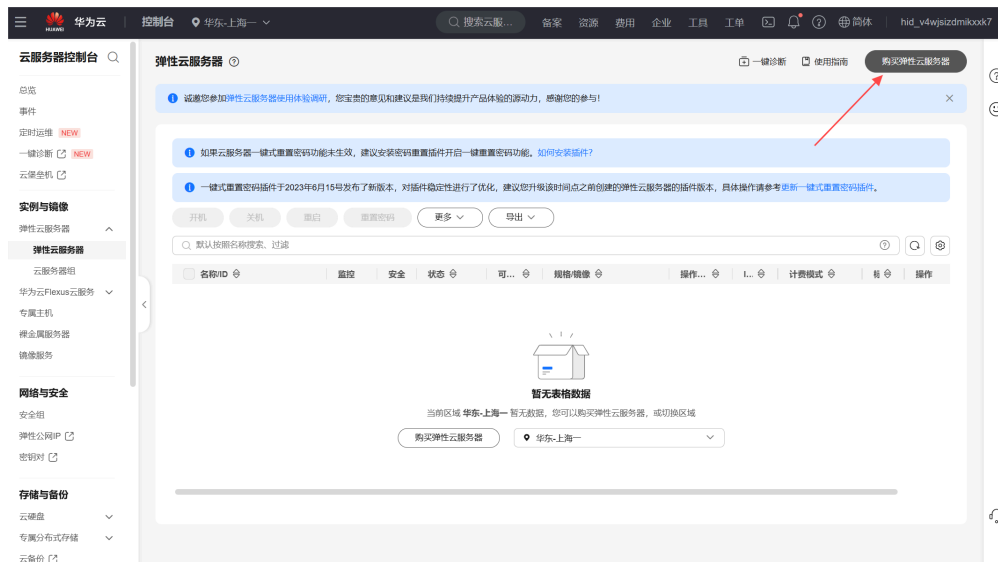
1. Development Environment Construction

1.1. Server Purchase

On the Huawei Cloud ECS homepage, search "ECS" and click "Management Console" to enter the ECS management page.



Create an elastic cloud server Select "North China-Beijing Four" in the console area, select "Elastic Cloud Server" in the left menu bar, and "Buy Elastic Cloud Server" in the upper right corner.



In the Basic Configuration, select the following configuration:

- Billing mode: billing on demand.
- Region: North China-Beijing 4th.
- Available area: Random allocation.
- CPU architecture: x86 calculation.
- Specification: General calculation enhanced type | c7.2xlarge.2 | 8vCPUs | 16 GiB
- Mirror: Public mirror, Ubuntu, Ubuntu 18.04 server 64bit
- System disk: Universal SSD, 100GB.



Click "Next: Network Configuration" in the next right corner of the window In the Network Configuration, select the following configuration:

- Network: You can go to the console to create a new virtual private cloud.
- Expand the network card: no.
- Security group: You can create a new security group.
- Flexible public network IP:
- Buy it now. Line: full dynamic BGP.
- Public network bandwidth: charge by traffic.
- Broadband size: custom, 200 Mbit/s.

- Release behavior: Check the Release with the instance.

网络

虚拟私有云 [?](#)

vpc-default(192.168.0.0/16) [新建虚拟私有云](#)

主网卡

subnet-default(192.168.0.0/24) 自动分配IP地址 可用私有IP数量250个

⊕ 新增扩展网卡

您还可以增加 3 块网卡

源/目的检查 [?](#)

☒

安全组

选择安全组 [?](#)

default(005c8838-f821-467b-8103-a89989169f49) [新建安全组](#)

请确保所选安全组已放通22端口 (Linux SSH登录) , 3389端口 (Windows远程登录) 和 ICMP 协议 (Ping) 。 [配置安全组规则](#)

展开安全组规则 [v](#)

公网访问

Click "Next: the advanced configuration" in the lower right corner of the window In the Advanced Configuration section, select the following configuration:

Cloud server name: It can be customized.

Login certificate: password.

User name: root.

Password: custom (subsequent login use, remember). C

loud backup: it is not purchased temporarily.

Cloud Server Group: None. Advanced option: None.

三 华为云 控制台

搜索云产品 备案 资源 费用 企业 工具 工单 帮助 简体中文 hid_v4wjszdmkoo

☒ 开启详细云监控 [免费](#) [?](#)

开启后云服务器CPU、内存、网络、磁盘、进程等指标的1分钟详细监控

☐ 云服务器组 [?](#)

通过云服务器组功能，弹性云服务器在创建时，将尽量分散地创建在不同的主机上，提高业务的可靠性。

云服务器描述

0/85

委托 [?](#)

--请选择-- [新建委托](#)

CPU 选项 [?](#)

☐ 指定 CPU 选项

购买量

使用时长

☐ 设置定时删除时间 [?](#)

购买数量

1

您最多可以创建200台云服务器。申请更多云服务器配额请单击[申请扩大配额](#)

配置概览

存储与备份

系统盘: 通用型SSD, 100GB, SCSI

网络

虚拟私有云: vpc-default(192.168.0.0/16)

主网卡: subnet-default(192.168.0.0/24)

源/目的检查: 开启

安全组

default

公网访问

弹性公网IP: 全动态BGP | 按流量计费 | 200 Mbits | 随实例释放

云服务器管理

云服务器名称: ecs-6715

登录凭证: 密码

标签: --

高级配置

云监控: 已开启详细监控 [免费](#)

购买量

定时删除时间: --

购买数量: 1

☒ 我已阅读并同意《弹性云服务器服务声明》[《弹性云服务器服务声明》](#) [《弹性云服务器服务声明》](#)

[《弹性云服务器服务声明》](#)

对外提供网站服务，需完成ICP备案。 [了解ICP备案](#)

立即购买

购买量: 1 台

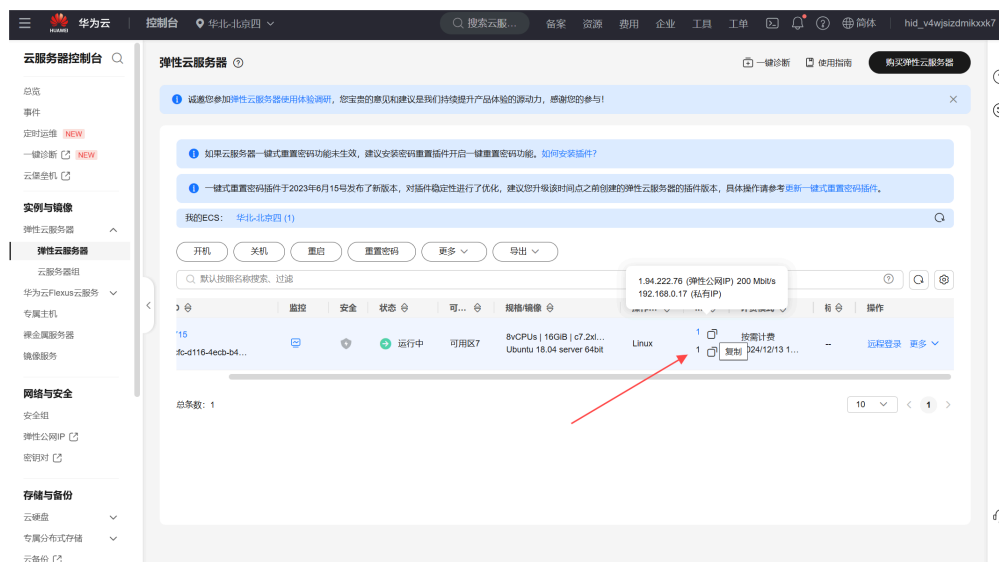
配置费用 ¥1.81 /小时 + 弹性公网IP流量费用 ¥0.80 /GB

实际扣费以账单为准。 [了解计费详情](#)

Click "Next: Confirm configuration" in the lower right corner of the window, In Confirm Configuration, select the following configuration: Agreement: Check that I have read and agreed to the Mirror Disclaimer.

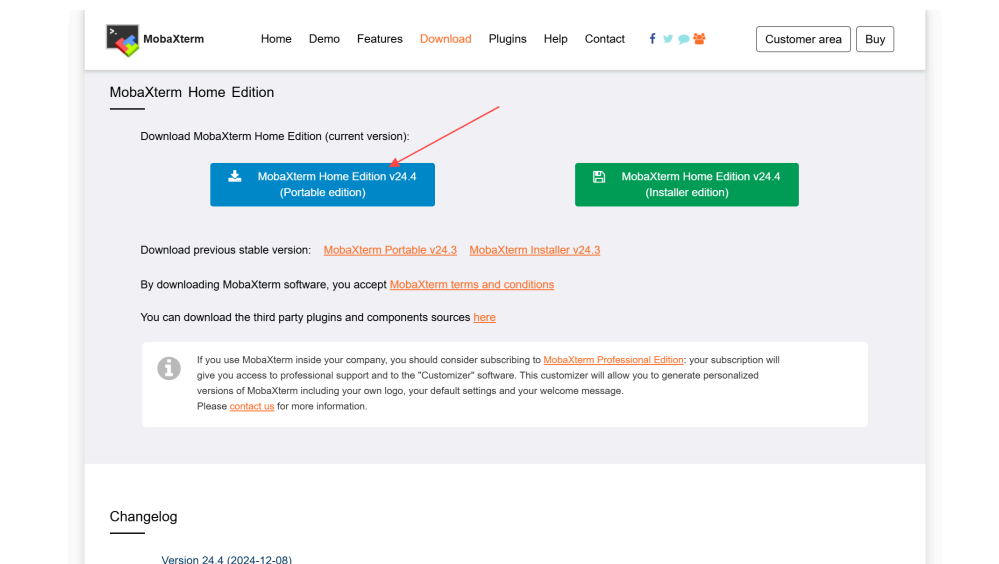
Click "Buy Now" in the lower right corner of the window. After the Task Submission succeeded, select Return to Server List to return to the management console of the elastic cloud server and see that the ECS created is running.

Note the flexible public network IP address displayed in the IP Address, which will be used later.



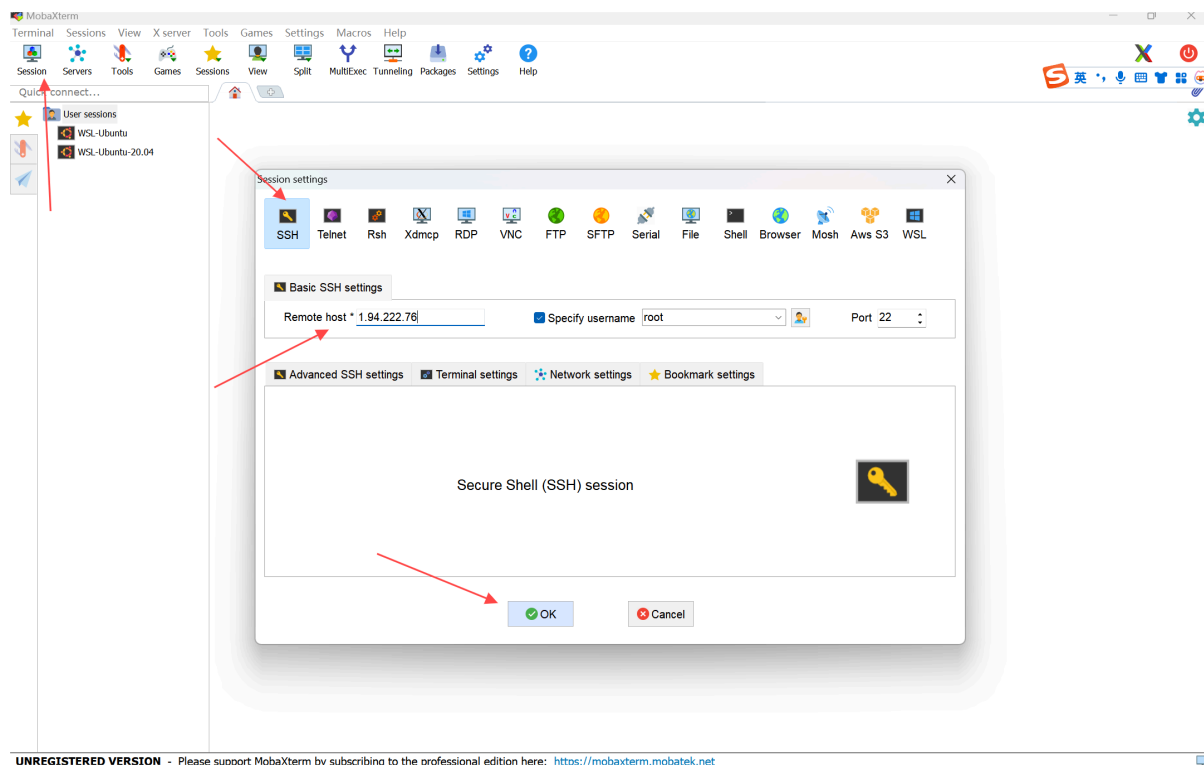
1.2. MobaXterm Connect ECS

Download the MobaXterm Go to MobaXterm's official website home page: <https://mobaxterm.mobatek.net/>
Select the Home Edition and download the MobaXterm Home Edition v21.x (Portable edition). Unpack the MobaXterm_Portable_v21.x.zip file after the download is complete.

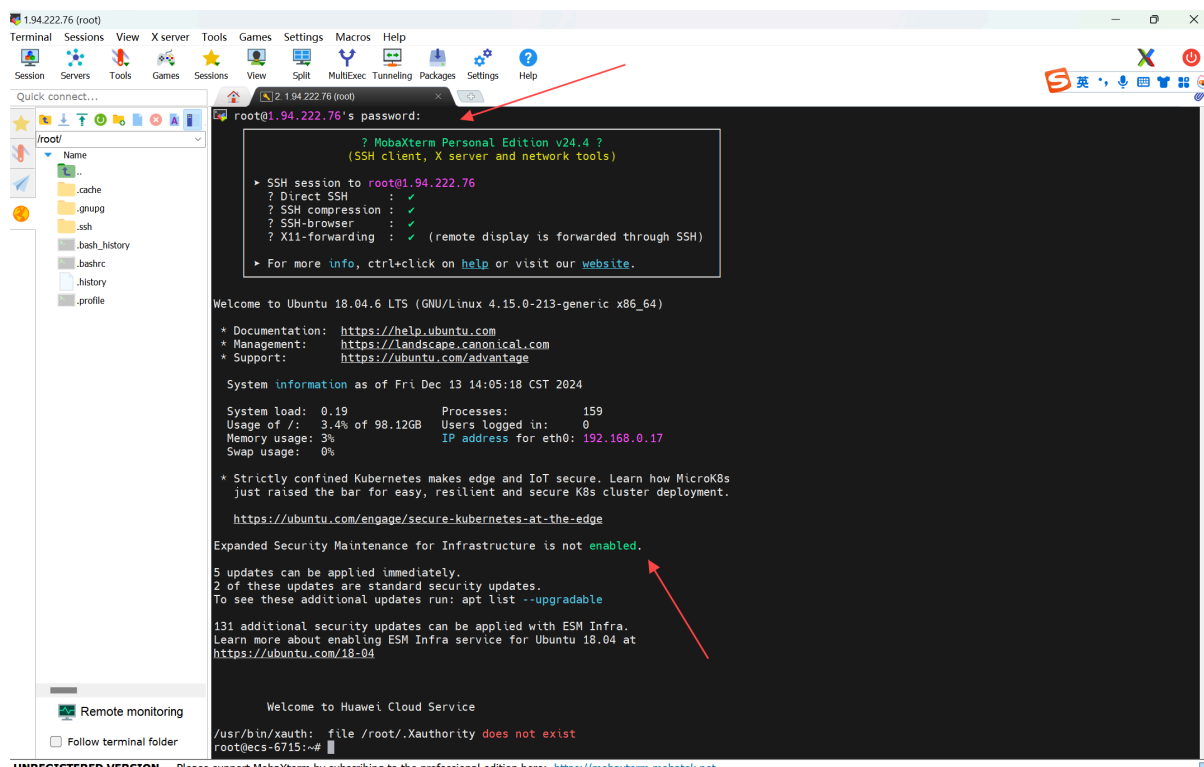


Use MobaXterm to connect to an elastic cloud server Enter the decompression MobaXterm_Portable_v21.x folder, Open the MobaXterm_Personal_21.x.exe file, Select the "Session" of the menu bar, Then enter the "Session settings" page, Remote link selection "SSH" protocol, Enter Figure 2-12 The elastic public network IP address displayed when the ECS elastic cloud server is created, Select the specified user name "Specify username", User name is "root", Select OK for submission after the configuration is complete.

名称	修改日期	类型	大小
CygUtils.plugin	2024/12/13 14:01	PLUGIN 文件	17,748 KB
CygUtils64.plugin	2024/12/13 14:01	PLUGIN 文件	11,723 KB
MobaXterm_Personal_24.4.exe	2024/12/13 14:01	应用程序	16,645 KB



MobaXterm Login to ECS requires a password. In step 4 of ECS elastic cloud server, the elastic cloud server root user password has been customized in the advanced configuration. You can enter here. MobaXterm The remote link to the elastic cloud server is successful, and the cloud environment of the elastic cloud server should be further configured later.



2. Code and Data Download

2.1. Get Code Files

Use git to download the source code of the training script from mindspore, switch to the home directory, create a working directory such as / work, and execute the following command:

```
git clone https://gitee.com/mindspore/models.git
```

The deepspeech2 project code for this experiment is located at models / research / audio / deepspeech2.

Training and inference-related parameters in the config.py file.

```
root@ecs-6715:~# cd "/home/"
root@ecs-6715:/home# mkdir work
root@ecs-6715:/home# cd work
root@ecs-6715:/home/work# ^C
root@ecs-6715:/home/work# git clone https://gitee.com/mindspore/models.git
Cloning into 'models'...
remote: Enumerating objects: 87074, done.
remote: Total 87074 (delta 0), reused 0 (delta 0), pack-reused 87074
Receiving objects: 100% (87074/87074), 502.99 MiB | 22.09 MiB/s, done.
Resolving deltas: 100% (62429/62429), done.
Checking out files: 100% (18471/18471), done.
root@ecs-6715:/home/work#
```

2.2. Dataset and its Preprocessing

2.2.1. Download the LibriSpeech Dataset

The link to download the data set is: <http://www.openslr.org/12>.

- **training set**

Train-clean-100: [6.3G] (100 Hour No Noise Speech Training Set) (just download this file)

- **validation set**

dev-clean.tar.gz [337M] (No Noise)

dev-other.tar.gz [314M] (with noise)

- **test set**

test-clean.tar.gz [346M] (Test set, No Noise)

test-other.tar.gz [328M] (Test set, noisy)

LibriSpeech Data directory structure, is as follows:

```
|—LibriSpeech
|— train
|   |—train-clean-100
|— val
|   |—dev-clean.tar.gz
|   |—dev-other.tar.gz
|— test_other
|   |—test-other.tar.gz
|—test_clean
|   |—test-clean.tar.gz
```

2.2.2. Install the Python3.9.0

Installing python dependency and software such as gcc.

```
sudo apt-get install -y gcc g++ make cmake zlib1g zlib1g-dev openssh libsqlite3-dev libssl-
dev libffi-dev unzip pciutils net-tools libblas-dev gfortran libblas3 libopenblas-dev
libgmp-dev sox libjpeg8-dev
```

```

root@ecs-6715:/home/Python-3.9.0# sudo apt-get install -y gcc g++ make cmake zlib1g zlib1g-dev openssl libsqlite3-dev libssl-dev libffi-dev unzip
pciutils net-tools libblas-dev gfortran libblas3 libopenblas-dev libgmp-dev sox libjpeg8-dev
Reading package lists... Done
Building dependency tree
Reading state information... Done
make is already the newest version (4.1-9.1ubuntu1).
make set to manually installed.
net-tools is already the newest version (1.60+git20161116.90da8a0-1ubuntu1).
g++ is already the newest version (4:7.4.0-1ubuntu2.3).
g++ set to manually installed.
gcc is already the newest version (4:7.4.0-1ubuntu2.3).
gcc set to manually installed.
openssl is already the newest version (1.1.1-1ubuntu2.1~18.04.23).
pciutils is already the newest version (1:3.5.2-1ubuntu1.1).
unzip is already the newest version (6.0-21ubuntu1.2).

```

Use wget to download the python3.9.0 source package, which can be downloaded to any directory of the installation environment with the command:

```

wget https://www.python.org/ftp/python/3.9.0/Python-3.9.0.tgz
tar -zxvf Python-3.9.0.tgz

```

```

root@ecs-6715:/home# wget https://www.python.org/ftp/python/3.9.0/Python-3.9.0.tgz
--2024-12-13 14:37:26-- https://www.python.org/ftp/python/3.9.0/Python-3.9.0.tgz
Resolving www.python.org (www.python.org)... 151.101.88.223, 2a04:4e42:15::223
Connecting to www.python.org (www.python.org)|151.101.88.223|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 26724009 (25M) [application/octet-stream]
Saving to: 'Python-3.9.0.tgz'

Python-3.9.0.tgz      100%[=====] 25.49M  142KB/s  in 3m 37s
2024-12-13 14:41:03 (120 KB/s) - 'Python-3.9.0.tgz' saved [26724009/26724009]

```

```

root@ecs-6715:/home# tar -zxvf Python-3.9.0.tgz
Python-3.9.0/
Python-3.9.0/CODE_OF_CONDUCT.md
Python-3.9.0/README.rst
Python-3.9.0/Doc/
Python-3.9.0/Doc/howto/
Python-3.9.0/Doc/howto/pyporting.rst
Python-3.9.0/Doc/howto/logging-cookbook.rst
Python-3.9.0/Doc/howto/logging_flow.png
Python-3.9.0/Doc/howto/sorting.rst
Python-3.9.0/Doc/howto/functional.rst
Python-3.9.0/Doc/howto/regex.rst
Python-3.9.0/Doc/howto/ipaddress.rst
Python-3.9.0/Doc/howto/argparse.rst
Python-3.9.0/Doc/howto/urllib2.rst
Python-3.9.0/Doc/howto/unicode.rst
Python-3.9.0/Doc/howto/index.rst
Python-3.9.0/Doc/howto/logging.rst
Python-3.9.0/Doc/howto/curses.rst
Python-3.9.0/Doc/howto/descriptor.rst
Python-3.9.0/Doc/howto/sockets.rst
Python-3.9.0/Doc/howto/instrumentation.rst
Python-3.9.0/Doc/howto/cporting.rst
Python-3.9.0/Doc/howto/clinic.rst
Python-3.9.0/Doc/runtime.txt
Python-3.9.0/Doc/README.rst
Python-3.9.0/Doc/install/

```

Go to the decompression folder and execute the configuration, compile, and installation commands:

```

cd Python-3.9.0
chmod +x configure
./configure --prefix=/usr/local/python3.9.0 --enable-shared
make
sudo make install

```

```

root@ecs-6715:/home/Python-3.9.0# chmod +x configure # configure文件添加可执行权限
root@ecs-6715:/home/Python-3.9.0# ./configure --prefix=/usr/local/python3.9.0 --enable-shared
checking build system type... x86_64-pc-linux-gnu
checking host system type... x86_64-pc-linux-gnu
checking for python3.9... no
checking for python3... python3
checking for --enable-universalsdk... no
checking for --with-universal-archs... no
checking MACHDEP... "linux"
checking for gcc... gcc
checking whether the C compiler works... yes
checking for C compiler default output file name... a.out
checking for suffix of executables...
checking whether we are cross compiling... no
checking for suffix of object files... o
checking whether we are using the GNU C compiler... yes
checking whether gcc accepts -g... yes
checking for gcc option to accept ISO C89... none needed
checking how to run the C preprocessor... gcc -E
checking for grep that handles long lines and -e... /bin/grep
checking for a sed that does not truncate output... /bin/sed
checking for --with-cxx-main=<compiler>... no
checking for g++... no
configure:

By default, distutils will build C++ extension modules with "g++".
If this is not intended, then set CXX on the configure command line.

checking for the platform triplet based on compiler characteristics... x86_64-linux-gnu
checking for -Wl,--no-as-needed... yes
checking for egrep... /bin/grep -E
checking for ANSI C header files... yes
checking for sys/types.h... yes

```

```

then rm -f /usr/local/python3.9.0/bin/python3; \
else true; \
fi
cd /usr/local/python3.9.0/bin; ln -s python3.9 python3)
if test "3.9" != "3.9"; then \
rm -f /usr/local/python3.9.0/bin/python3.9-config; \
(cd /usr/local/python3.9.0/bin; ln -s python3.9-config python3.9-config); \
rm -f /usr/local/python3.9.0/lib/pkgconfig/python-3.9.pc; \
(cd /usr/local/python3.9.0/lib/pkgconfig; ln -s python-3.9.pc python-3.9.pc); \
rm -f /usr/local/python3.9.0/lib/pkgconfig/python-3.9-embed.pc; \
(cd /usr/local/python3.9.0/lib/pkgconfig; ln -s python-3.9-embed.pc python-3.9-embed.pc); \
fi
rm -f /usr/local/python3.9.0/bin/python3-config
(cd /usr/local/python3.9.0/bin; ln -s python3.9-config python3.9-config)
rm -f /usr/local/python3.9.0/lib/pkgconfig/python3.pc
cd /usr/local/python3.9.0/lib/pkgconfig; ln -s python-3.9.pc python3.pc)
rm -f /usr/local/python3.9.0/lib/pkgconfig/python3-embed.pc
cd /usr/local/python3.9.0/lib/pkgconfig; ln -s python-3.9-embed.pc python3-embed.pc)
rm -f /usr/local/python3.9.0/bin/idle3
(cd /usr/local/python3.9.0/bin; ln -s idle3.9 idle3)
rm -f /usr/local/python3.9.0/bin/pydoc3
cd /usr/local/python3.9.0/bin; ln -s pydoc3.9 pydoc3)
rm -f /usr/local/python3.9.0/bin/2to3
cd /usr/local/python3.9.0/bin; ln -s 2to3-3.9 2to3)
if test "x" != "x" ; then \
rm -f /usr/local/python3.9.0/bin/python3-32; \
(cd /usr/local/python3.9.0/bin; ln -s python3.9-32 python3-32) \
fi
rm -f /usr/local/python3.9.0/share/man/man1/python3.1
cd /usr/local/python3.9.0/share/man/man1; ln -s python3.9.1 python3.1)
if test "xupgrade" != "xno" ; then \
case upgrade in \
upgrade) ensurepip="--upgrade" ;; \
install|*) ensurepip="" ;; \
esac; \
LD_LIBRARY_PATH=/home/Python-3.9.0 ./python -E -m ensurepip \
$ensurepip --root=/ ; \
fi
Looking in links: /tmp/tmp15ynm6
Processing /tmp/tmp15ynm6/setuptools-49.2.1-py3-none-any.whl
Processing /tmp/tmp15ynm6/pip-20.2.3-py2.py3-none-any.whl
Installing collected packages: setuptools, pip
WARNING: The script easy_install-3.9 is installed in '/usr/local/python3.9.0/bin' which is not on PATH.
Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
WARNING: The scripts pip3 and pip3.9 are installed in '/usr/local/python3.9.0/bin' which is not on PATH.
Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
Successfully installed pip-20.2.3 setuptools-49.2.1
root@ecs-6715:/home/Python-3.9.0#

```

Query whether there is libpython3.9.so.1.0 under /usr/lib64 or /usr/lib, skip this step or back up the libpython3.9.so.1.0 file with the following command.

```
cp /usr/local/python3.9.0/lib/libpython3.9.so.1.0 /usr/lib
```

```

root@ecs-6715:/home/Python-3.9.0# cp /usr/local/python3.9.0/lib/libpython3.9.so.1.0 /usr/lib
root@ecs-6715:/home/Python-3.9.0#

```

Perute the following command setting soft link:

```

sudo ln -s /usr/local/python3.9.0/bin/python3.9 /usr/bin/python
sudo ln -s /usr/local/python3.9.0/bin/pip3.9 /usr/bin/pip
sudo ln -s /usr/local/python3.9.0/bin/python3.9 /usr/bin/python3.9
sudo ln -s /usr/local/python3.9.0/bin/pip3.9 /usr/bin/pip3.9

```



```
root@ecs-6715:/home/Python-3.9.0# sudo ln -s /usr/local/python3.9.0/bin/python3.9 /usr/bin/python
root@ecs-6715:/home/Python-3.9.0# sudo ln -s /usr/local/python3.9.0/bin/pip3.9 /usr/bin/pip
root@ecs-6715:/home/Python-3.9.0# sudo ln -s /usr/local/python3.9.0/bin/python3.9 /usr/bin/python3.9
root@ecs-6715:/home/Python-3.9.0# sudo ln -s /usr/local/python3.9.0/bin/pip3.9 /usr/bin/pip3.9
```

After the installation is complete, perform the following command to view the installation version, and if the relevant version information is returned, the installation is successful.

```
python3.9 --version
pip3.9 --version
```

```
root@ecs-6715:/home/Python-3.9.0# python3.9 --version
Python 3.9.0
root@ecs-6715:/home/Python-3.9.0# pip3.9 --version
pip 20.2.3 from /usr/local/python3.9.0/lib/python3.9/site-packages/pip (python 3.9)
root@ecs-6715:/home/Python-3.9.0#
```

2.2.3. Install the MindSpore and the Required Dependency Package

Install mindspore, install according to the actual server architecture, please refer to <https://www.mindspore.cn/install>.

```
pip install https://ms-release.obs.cn-north-4.myhuaweicloud.com/1.6.0/MindSpore/cpu/x86_64/mindspore-1.6.0-cp39-cp39-linux_x86_64.whl \
--trusted-host ms-release.obs.cn-north-4.myhuaweicloud.com \
-i https://pypi.tuna.tsinghua.edu.cn/simple
```

```
root@ecs-6715:/home/Python-3.9.0# pip install https://ms-release.obs.cn-north-4.myhuaweicloud.com/1.6.0/MindSpore/cpu/x86_64/mindspore-1.6.0-cp39-cp39-linux_x86_64.whl \
> --trusted-host ms-release.obs.cn-north-4.myhuaweicloud.com \
> -i https://pypi.tuna.tsinghua.edu.cn/simple
Looking in indexes: https://pypi.tuna.tsinghua.edu.cn/simple
Collecting mindspore==1.6.0
  Downloading https://ms-release.obs.cn-north-4.myhuaweicloud.com/1.6.0/MindSpore/cpu/x86_64/mindspore-1.6.0-cp39-cp39-linux_x86_64.whl (95.7 MB)
    | 95.7 MB 218 kB/s
Collecting numpy>=1.17.0
  Downloading https://pypi.tuna.tsinghua.edu.cn/packages/b9/14/78635daab4b07c0930c919d451b8bf8c164774e6a3413aed04a6d95758ce/numpy-2.0.2-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (19.5 MB)
    | 19.5 MB 22.8 MB/s
Collecting psutil>=5.6.1
  Downloading https://pypi.tuna.tsinghua.edu.cn/packages/58/4d/8245e6f76a93c98aab285a43ea71ff1b171bcd90c9d238bf81f7021fb233/psutil-6.1.0-cp36-abi3-manylinux_2_12_x86_64.manylinux2010_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_64.whl (287 kB)
    | 287 kB 5.2 MB/s
Collecting protobuf>=3.13.0
  Downloading https://pypi.tuna.tsinghua.edu.cn/packages/04/52/c97c58a33b3d6c89a8138788576d372a90a6556f354799971c6b4d16d871/protobuf-5.29.1-cp38-abi3-manylinux2014_x86_64.whl (319 kB)
    | 319 kB 116.5 MB/s
Collecting packaging>=20.0
  Downloading https://pypi.tuna.tsinghua.edu.cn/packages/88/ef/eb23f262cca3c0c4eb7ab1933c3b1f03d021f2c48f54763065b6f0e321be/packaging-24.2-py3-none-any.whl (65 kB)
    | 65 kB 3.6 MB/s
Collecting pillow>=6.2.0
  Downloading https://pypi.tuna.tsinghua.edu.cn/packages/14/b1/c8f428bae932a27ce9c87e7b21aba8ea3e820aa11413c5a795868c37e039/pillow-11.0.0-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (4.3 MB)
    | 4.3 MB 10.1 MB/s
Collecting asttokens>=2.0.0
  Downloading https://pypi.tuna.tsinghua.edu.cn/packages/25/8a/c46dc25341b5bce5472c718902eb3d38600a903b14fa6aececf3f21a6f/asttokens-3.0.0-py3-none-any.whl (26 kB)
Collecting scipy>=1.5.2
  Downloading https://pypi.tuna.tsinghua.edu.cn/packages/35/f5/d0ad1a96f80962ba65e2ce1de6a1e59edecdf0a7b55990ed208848012e0/scipy-1.13.1-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (38.6 MB)
    | 38.6 MB 3.2 MB/s
Installing collected packages: numpy, psutil, protobuf, packaging, pillow, asttokens, scipy, mindspore
Successfully installed asttokens-3.0.0 mindspore-1.6.0 numpy-2.0.2 packaging-24.2 pillow-11.0.0 protobuf-5.29.1 psutil-6.1.0 scipy-1.13.1
WARNING: You are using pip version 20.2.3; however, version 24.3.1 is available.
You should consider upgrading via the '/usr/local/python3.9.0/bin/python3.9 -m pip install --upgrade pip' command.
```

Pip source installation, you can add a mirror source installation when the dependent package file is large, such as pip install-i <https://pypi.tuna.tsinghua.edu.cn/simple> sox.

```
pip3.9 install wget
pip3.9 install tqdm
pip3.9 install sox
```

```
root@ecs-6715:/home/Python-3.9.0# pip3.9 install wget
Collecting wget
  Downloading wget-3.2.zip (10 kB)
Using legacy 'setup.py install' for wget, since package 'wheel' is not installed.
Installing collected packages: wget
  Running setup.py install for wget ... done
Successfully installed wget-3.2
WARNING: You are using pip version 20.2.3; however, version 24.3.1 is available.
You should consider upgrading via the '/usr/local/python3.9.0/bin/python3.9 -m pip install --upgrade pip' command.
root@ecs-6715:/home/Python-3.9.0# pip3.9 install tqdm
Collecting tqdm
  Downloading tqdm-4.67.1-py3-none-any.whl (78 kB)
    | 78 kB 620 kB/s
Installing collected packages: tqdm
Successfully installed tqdm-4.67.1
WARNING: You are using pip version 20.2.3; however, version 24.3.1 is available.
You should consider upgrading via the '/usr/local/python3.9.0/bin/python3.9 -m pip install --upgrade pip' command.
```

```

root@ecs-6715:/home/Python-3.9.0# pip install typing-extensions
Collecting typing-extensions
  Downloading typing_extensions-4.12.2-py3-none-any.whl (37 kB)
Installing collected packages: typing-extensions
Successfully installed typing-extensions-4.12.2
WARNING: You are using pip version 20.2.3; however, version 24.3.1 is available.
You should consider upgrading via the '/usr/local/python3.9.0/bin/python3.9 -m pip install --upgrade pip' command.
root@ecs-6715:/home/Python-3.9.0# pip3.9 install sox
Collecting sox
  Using cached sox-1.5.0.tar.gz (63 kB)
Requirement already satisfied: numpy>=1.9.0 in /usr/local/python3.9.0/lib/python3.9/site-packages (from sox) (2.0.2)
Requirement already satisfied: typing-extensions>=3.7.4.2 in /usr/local/python3.9.0/lib/python3.9/site-packages (from sox) (4.12.2)
Using legacy 'setup.py install' for sox, since package 'wheel' is not installed.
Installing collected packages: sox
  Running setup.py install for sox ... done
Successfully installed sox-1.5.0
WARNING: You are using pip version 20.2.3; however, version 24.3.1 is available.
You should consider upgrading via the '/usr/local/python3.9.0/bin/python3.9 -m pip install --upgrade pip' command.

```

2.2.4. The Data Preprocessing SeanNaren Scripts was Downloaded

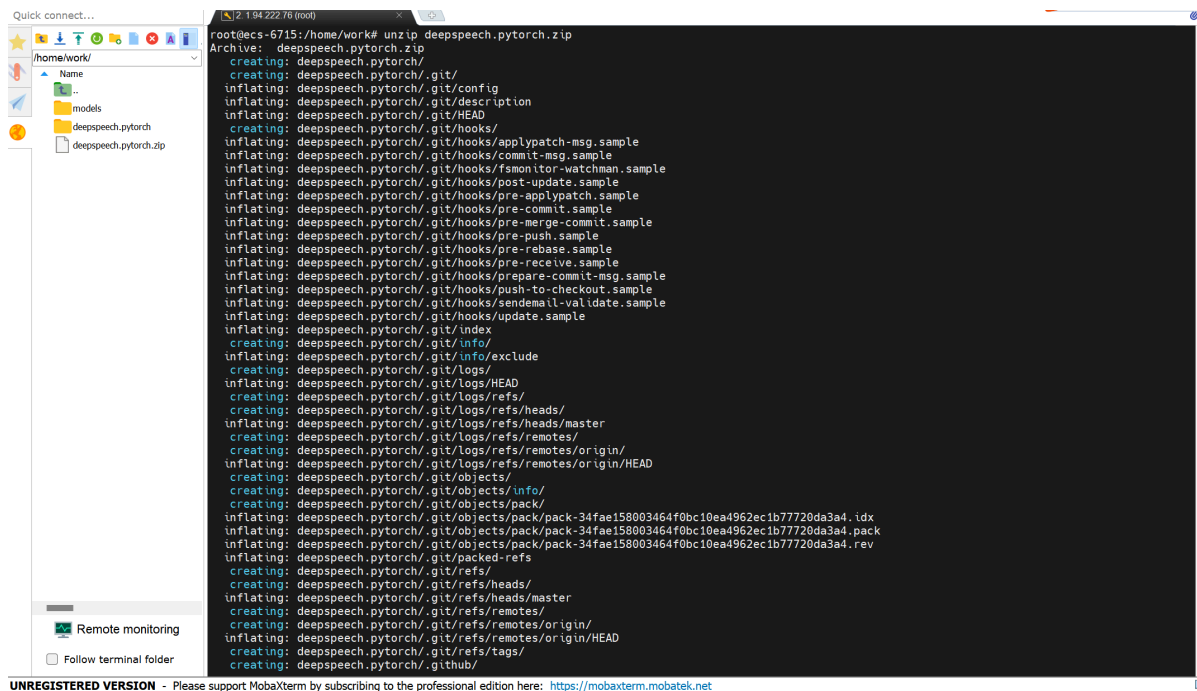
After MobaXterm / Finalshell (recommended) connects to the ECS server, switch to the home directory, create the working directory, and then use the scripts in SeanNaren to process the data. SeanNaren Script link: <http://github.com/SeanNaren/deepspeech.pytorch>.

```

cd ../home
mkdir work
cd work
git clone https://github.com/SeanNaren/deepspeech.pytorch.git

```

Here, you need to access the external network, it is recommended to download directly to the local and then upload the zip file before decompression



2.2.5. LibriSpeech Data Preprocessing

The training set of train-clean-100 was downloaded locally via the dataset link <http://www.openslr.org/12>, validation sets dev-clean.tar.gz and dev-other.tar.gz, and test sets test-clean.tar.gz and test-other.tar.gz. Upload the local data set to the MobaXterm server.

The structure of the data set is shown below:

```

root@ecs-6715:/home/work/deepspeech.pytorch# tree LibriSpeech_dataset
LibriSpeech_dataset
├── test_clean
│   └── test-clean.tar.gz
├── test_other
│   └── test-other.tar.gz
├── train
│   └── train-clean-100.tar.gz
└── val
    ├── dev-clean.tar.gz
    └── dev-other.tar.gz
4 directories, 5 files

```

Copy the librispeech.py in the data directory of the deepspeech.pytorch to the deepspeech.pytorch directory and execute the following command:

```

cd deepspeech.pytorch
cp ./data/librispeech.py ./

```

```

root@ecs-6715:/home/work# cd deepspeech.pytorch
root@ecs-6715:/home/work/deepspeech.pytorch# cp ./data/librispeech.py ./

```

Modify the librispeech.py code data set path, refer to step 3, set the directory structure in the current directory, and change the code path to the data set actual path, as shown in the figure below:

```

LIBRI_SPEECH_URLS = {
    "train": ["LibriSpeech_dataset/train/train-clean-100.tar.gz"],

    "val": ["LibriSpeech_dataset/val/dev-clean.tar.gz",
            "LibriSpeech_dataset/val/dev-other.tar.gz"],

    "test_clean": ["LibriSpeech_dataset/test-clean/test-clean.tar.gz"],
    "test_other": ["LibriSpeech_dataset/test-other/test-other.tar.gz"]
}

```

Execute the data set processing command, and execute the command as follows.

```
python librispeech.py
```

```

root@ecs-6715:/home/work/deepspeech.pytorch# python librispeech.py
Unpacking train-clean-100.tar.gz...
Converting flac files to wav and extracting transcripts...
838it [03:53, 3.59it/s]
Finished LibriSpeech_dataset/train/train-clean-100.tar.gz
Gathering durations...
100% | 28539/28539 [00:23<00:00, 1197.48it/s]
Sorting manifests...
Pruning manifests between 1 and 15 seconds
Total duration of split: 233513.3018s
100% | 20336/20336 [00:00<00:00, 54626.48it/s]
Unpacking dev-clean.tar.gz...
Converting flac files to wav and extracting transcripts...
139it [00:16, 8.35it/s]
Finished LibriSpeech_dataset/val/dev-clean.tar.gz
Unpacking dev-other.tar.gz...
Converting flac files to wav and extracting transcripts...
126it [00:16, 7.47it/s]
Finished LibriSpeech_dataset/val/dev-other.tar.gz
Gathering durations...
100% | 5567/5567 [00:03<00:00, 1398.12it/s]
Sorting manifests...
Total duration of split: 37832.3861s
100% | 5567/5567 [00:00<00:00, 58691.16it/s]
Unpacking test-clean.tar.gz...
Converting flac files to wav and extracting transcripts...
129it [00:15, 8.17it/s]
Finished LibriSpeech_dataset/test_clean/test-clean.tar.gz
Gathering durations...
100% | 2620/2620 [00:01<00:00, 1419.30it/s]
Sorting manifests...
Total duration of split: 19452.4806s
100% | 2620/2620 [00:00<00:00, 56671.31it/s]
Unpacking test-other.tar.gz...
Converting flac files to wav and extracting transcripts...
125it [00:16, 7.37it/s]
Finished LibriSpeech_dataset/test_other/test-other.tar.gz
Gathering durations...
100% | 2939/2939 [00:02<00:00, 1300.47it/s]
Sorting manifests...
Total duration of split: 19229.5701s
100% | 2939/2939 [00:00<00:00, 58223.68it/s]
root@ecs-6715:/home/work/deepspeech.pytorch#

```

After the data processing, the data directory structure is as follows:

```

├─ LibriSpeech_dataset
│   ├── train
│   │   ├── wav
│   │   └── txt

```

```

|   ├── val
|   |   ├── wav
|   |   └── txt
|   ├── test_clean
|   |   ├── wav
|   |   └── txt
|   └── test_other
|       ├── wav
|       └── txt
└── libri_test_clean_manifest.json, libri_test_other_manifest.json,
    libri_train_manifest.json, libri_val_manifest.json

```

Go from the json file to the csv file, create the json_to_csv.py in the deepspeech.pytorch directory, and copy the code to the file with the following code:

```
touch json_to_csv.py
```

```

# json_to_csv.py:
import json
import csv
import argparse

parser = argparse.ArgumentParser(description='Image classification')
parser.add_argument("--json", type=str, default="", help="")
parser.add_argument("--csv", type=str, default="", help="")
config = parser.parse_args()

def trans(jsonfile, csvfile):
    jsonData = open(jsonfile)
    csvfile = open(csvfile, "a")
    for i in jsonData:
        dic = json.loads(i[0:])
        root_path = dic["root_path"]
        for j in dic["samples"]:
            wav_path = j["wav_path"]
            transcript_path = j["transcript_path"]
            res_wav = root_path + '/' + wav_path
            res_txt = root_path + '/' + transcript_path
            res = [res_wav, res_txt]
            writer = csv.writer(csvfile)
            writer.writerow(res)
    jsonData.close()
    csvfile.close()

if __name__ == "__main__":
    trans(config.json, config.csv)

```

Run the command as shown in the following figure:

```
python json_to_csv.py --json libri_test_clean_manifest.json --csv
libri_test_clean_manifest.csv
python json_to_csv.py --json libri_test_other_manifest.json --csv
libri_test_other_manifest.csv
python json_to_csv.py --json libri_train_manifest.json --csv libri_train_manifest.csv
python json_to_csv.py --json libri_val_manifest.json --csv libri_val_manifest.csv
```

```
root@ecs-6715:/home/work/deepspeech.pytorch# touch json_to_csv.py
root@ecs-6715:/home/work/deepspeech.pytorch# python json_to_csv.py --json libri_test_clean_manifest.json --csv libri_test_clean_manifest.csv
python json_to_csv.py --json libri_test_other_manifest.json --csv libri_test_other_manifest.csv
python json_to_csv.py --json libri_train_manifest.json --csv libri_train_manifest.csv
python json_to_csv.py --json libri_val_manifest.json --csv libri_val_manifest.csvroot@ecs-6715:/home/work/deepspeech.pytorch# python json_to_csv.p
y --json libri_test_other_manifest.json --csv libri_test_other_manifest.csv
root@ecs-6715:/home/work/deepspeech.pytorch# python json_to_csv.py --json libri_train_manifest.json --csv libri_train_manifest.csv
root@ecs-6715:/home/work/deepspeech.pytorch# python json_to_csv.py --json libri_val_manifest.json --csv libri_val_manifest.csv
root@ecs-6715:/home/work/deepspeech.pytorch#
```

3. Model Training Results and Evaluation

3.1. Model Training

Switch to the models / official / audio / DeepSpeech2 directory, Model training requires the creation of the deepspeech_pytorch directory under the DeepSpeech2 directory and the decoder.py file under the deepspeech_pytorch directory.

```
mkdir deepspeech_pytorch
cd deepspeech_pytorch
touch decoder.py
```

Copy the code to the decoder.py file with the following code:

```
#!/usr/bin/env python
# -----
# Copyright 2015-2016 Nervana Systems Inc.
# Licensed under the Apache License, Version 2.0 (the "License");
# you may not use this file except in compliance with the License.
# You may obtain a copy of the License at
#
#     http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
# -----
# Modified to support pytorch Tensors

import Levenshtein as Lev
import torch
from six.moves import xrange

class Decoder(object):
    """
    Basic decoder class from which all other decoders inherit. Implements several
    helper functions. Subclasses should implement the decode() method.

    Arguments:
```

```

    labels (list): mapping from integers to characters.
    blank_index (int, optional): index for the blank '_' character. Defaults to 0.
    """

def __init__(self, labels, blank_index=0):
    self.labels = labels
    self.int_to_char = dict([(i, c) for (i, c) in enumerate(labels)])
    self.blank_index = blank_index
    space_index = len(labels) # To prevent errors in decode, we add an out of bounds
index for the space
    if ' ' in labels:
        space_index = labels.index(' ')
    self.space_index = space_index

def wer(self, s1, s2):
    """
    Computes the Word Error Rate, defined as the edit distance between the
    two provided sentences after tokenizing to words.
    Arguments:
        s1 (string): space-separated sentence
        s2 (string): space-separated sentence
    """

    # build mapping of words to integers
    b = set(s1.split() + s2.split())
    word2char = dict(zip(b, range(len(b))))

    # map the words to a char array (Levenshtein packages only accepts
    # strings)
    w1 = [chr(word2char[w]) for w in s1.split()]
    w2 = [chr(word2char[w]) for w in s2.split()]

    return Lev.distance(''.join(w1), ''.join(w2))

def cer(self, s1, s2):
    """
    Computes the Character Error Rate, defined as the edit distance.

    Arguments:
        s1 (string): space-separated sentence
        s2 (string): space-separated sentence
    """

    s1, s2, = s1.replace(' ', ''), s2.replace(' ', '')
    return Lev.distance(s1, s2)

def decode(self, probs, sizes=None):
    """
    Given a matrix of character probabilities, returns the decoder's
    best guess of the transcription

    Arguments:
        probs: Tensor of character probabilities, where probs[c,t]
                is the probability of character c at time t
        sizes(optional): Size of each sequence in the mini-batch
    Returns:
        string: sequence of the model's best guess for the transcription
    """

```

```

        """
        raise NotImplementedError

class BeamCTCDecoder(Decoder):
    def __init__(self,
                  labels,
                  lm_path=None,
                  alpha=0,
                  beta=0,
                  cutoff_top_n=40,
                  cutoff_prob=1.0,
                  beam_width=100,
                  num_processes=4,
                  blank_index=0):
        super(BeamCTCDecoder, self).__init__(labels)
        try:
            from ctctdecode import CTCBeamDecoder
        except ImportError:
            raise ImportError("BeamCTCDecoder requires paddleddecoder package.")
        labels = list(labels) # Ensure labels are a list before passing to decoder
        self._decoder = CTCBeamDecoder(labels, lm_path, alpha, beta, cutoff_top_n,
                                        cutoff_prob, beam_width,
                                        num_processes, blank_index)

    def convert_to_strings(self, out, seq_len):
        results = []
        for b, batch in enumerate(out):
            utterances = []
            for p, utt in enumerate(batch):
                size = seq_len[b][p]
                if size > 0:
                    transcript = ''.join(map(lambda x: self.int_to_char[x.item()],
                                        utt[0:size]))
                else:
                    transcript = ''
            utterances.append(transcript)
            results.append(utterances)
        return results

    def convert_tensor(self, offsets, sizes):
        results = []
        for b, batch in enumerate(offsets):
            utterances = []
            for p, utt in enumerate(batch):
                size = sizes[b][p]
                if sizes[b][p] > 0:
                    utterances.append(utt[0:size])
                else:
                    utterances.append(torch.tensor([], dtype=torch.int))
            results.append(utterances)
        return results

    def decode(self, probs, sizes=None):
        """

```

Decodes probability output using ctcdecode package.

Arguments:

probs: Tensor of character probabilities, where probs[c,t]
is the probability of character c at time t
sizes: Size of each sequence in the mini-batch

Returns:

string: sequences of the model's best guess for the transcription

"""

probs = probs.cpu()

out, scores, offsets, seq_lens = self._decoder.decode(probs, sizes)

strings = self.convert_to_strings(out, seq_lens)

offsets = self.convert_tensor(offsets, seq_lens)

return strings, offsets

```
class GreedyDecoder(Decoder):
```

```
    def __init__(self, labels, blank_index=0):
```

```
        super(GreedyDecoder, self).__init__(labels, blank_index)
```

```
    def convert_to_strings(self,
                           sequences,
                           sizes=None,
                           remove_repetitions=False,
                           return_offsets=False):
```

```
        """Given a list of numeric sequences, returns the corresponding strings"""
```

```
        strings = []
```

```
        offsets = [] if return_offsets else None
```

```
        for x in xrange(len(sequences)):
```

```
            seq_len = sizes[x] if sizes is not None else len(sequences[x])
```

```
            string, string_offsets = self.process_string(sequences[x], seq_len,
remove_repetitions)
```

```
            strings.append([string]) # we only return one path
```

```
            if return_offsets:
```

```
                offsets.append([string_offsets])
```

```
        if return_offsets:
```

```
            return strings, offsets
```

```
        else:
```

```
            return strings
```

```
    def process_string(self,
```

```
                      sequence,
```

```
                      size,
```

```
                      remove_repetitions=False):
```

```
        string = ''
```

```
        offsets = []
```

```
        for i in range(size):
```

```
            char = self.int_to_char[sequence[i].item()]
```

```
            if char != self.int_to_char[self.blank_index]:
```

```
                # if this char is a repetition and remove_repetitions=true, then skip
```

```
                if remove_repetitions and i != 0 and char == self.int_to_char[sequence[i -
1].item()]:
```

```
                    pass
```

```
                elif char == self.labels[self.space_index]:
```

```
                    string += ' '
```

```
                    offsets.append(i)
```



```

        else:
            string = string + char
            offsets.append(i)
        return string, torch.tensor(offsets, dtype=torch.int)

def decode(self, probs, sizes=None):
    """
    Returns the argmax decoding given the probability matrix. Removes
    repeated elements in the sequence, as well as blanks.

    Arguments:
        probs: Tensor of character probabilities from the network. Expected shape of
        batch x seq_length x output_dim
        sizes(optional): Size of each sequence in the mini-batch
    Returns:
        strings: sequences of the model's best guess for the transcription on inputs
        offsets: time step per character predicted
    """
    _, max_probs = torch.max(probs, 2)
    strings, offsets = self.convert_to_strings(max_probs.view(max_probs.size(0),
max_probs.size(1)),
                                           sizes,
                                           remove_repetitions=True,
                                           return_offsets=True)

    return strings, offsets

```

Model configuration Modify the config.py under the src. After the modification, "ctrl + s" is saved and exits.

Modify batch_size to 1 (one-process data size that is related to server device performance).

Modify epochs to 1 (about 48h, can be adjusted according to the actual demand).

Modifies the train_manifest to the libri_train_manifest.csv actual path.

Modifies test_manifest to libri_test_clean_manifest.csv to the actual path.

Modify the window type of eval_config to change hanning to hann.

Install the model python dependency:

```

cd /home/work/models/official/audio/DeepSpeech2
pip3.9 install -r requirements.txt
pip3.9 install Levenshtein
pip3.9 install -i https://pypi.tuna.tsinghua.edu.cn/simple torch==1.7.1
pip3.9 install numpy==1.20.0
pip install numba==0.53.1

```

```

root@ecs-6715:/home/work/models/official/audio/DeepSpeech2# pip3.9 install -r requirements.txt
Requirement already satisfied: numpy in /usr/local/python3.9.0/lib/python3.9/site-packages (from -r requirements.txt (line 1)) (2.0.2)
Collecting easydict
  Downloading easydict-1.13-py3-none-any.whl (6.8 kB)
Collecting librosa
  Downloading librosa-0.10.2.post1-py3-none-any.whl (260 kB)
    |#####| 260 kB 744 kB/s
Collecting soundfile
  Downloading soundfile-0.12.1-py2.py3-none-any.whl (24 kB)
Collecting audioread>=2.1.9
  Downloading audioread-3.0.1-py3-none-any.whl (23 kB)
Collecting msgpack>=1.0
  Downloading msgpack-1.1.0-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (377 kB)
    |#####| 377 kB 242 kB/s
Collecting scikit-learn>=0.20.0
  Downloading scikit_learn-1.6.0-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (13.5 MB)
    |#####| 4.6 MB 250 kB/s eta 0:00:36

```

Download the pre-training model, the download link is <https://ascend-professional-construction-dataset.obs.cn-north-4.myhuaweicloud.com/ASR/DeepSpeech.ckpt>, and the download command is as follows:

```
wget https://ascend-professional-construction-dataset.obs.cn-north-4.myhuaweicloud.com/ASR/DeepSpeech.ckpt
```

```
root@ecs-6715:/home/work/models/official/audio/DeepSpeech2# wget https://ascend-professional-construction-dataset.obs.cn-north-4.myhuaweicloud.com/ASR/DeepSpeech.ckpt
--2024-12-13 17:43:37-- https://ascend-professional-construction-dataset.obs.cn-north-4.myhuaweicloud.com/ASR/DeepSpeech.ckpt
Resolving ascend-professional-construction-dataset.obs.cn-north-4.myhuaweicloud.com (ascend-professional-construction-dataset.obs.cn-north-4.myhuaweicloud.com)... 100.125.83.133, 100.125.83.5, 100.125.76.5
Connecting to ascend-professional-construction-dataset.obs.cn-north-4.myhuaweicloud.com (ascend-professional-construction-dataset.obs.cn-north-4.myhuaweicloud.com)|100.125.83.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1039306262 (991M) [binary/octet-stream]
Saving to: 'DeepSpeech.ckpt'

DeepSpeech.ckpt      100%[=====] 991.16M  239MB/s   in 4.4s
2024-12-13 17:43:41 (227 MB/s) - 'DeepSpeech.ckpt' saved [1039306262/1039306262]
```

Modify the run_standalone_train_cpu.sh in the scripts directory to load the pre-training model The modification is as follows:

```
PATH_CHECKPOINT=$1
python ./train.py --device_target 'CPU' --pre_trained_model_path $PATH_CHECKPOINT

PATH_CHECKPOINT=$1
python ./train.py --device_target 'CPU' --pre_trained_model_path $PATH_CHECKPOINT
```

Train the model in the DeepSpeech2 directory and enter the following command.

```
bash scripts/run_standalone_train_cpu.sh PATH_CHECKPOINT
# PATH_CHECKPOINT: Pre-training file path
```

```
root@ecs-6715:/home/work/models/official/audio/DeepSpeech2# bash scripts/run_standalone_train_cpu.sh /home/work/models/official/audio/DeepSpeech2/DeepSpeech.ckpt
```

So that, the model is now being trained.

3.2. Model Training Results and Evaluation

3.2.1. View the Training Log

If we want to view the training log, the current directory under the train.log. Enter the command as follows:

```
tail -f train.log
```

```
root@ecs-6715:/home/work/models/official/audio# cd DeepSpeech2
root@ecs-6715:/home/work/models/official/audio/DeepSpeech2# bash scripts/run_standalone_train_cpu.sh /home/work/models/official/audio/DeepSpeech2/DeepSpeech.ckpt
root@ecs-6715:/home/work/models/official/audio/DeepSpeech2# tail -f train.log
epoch: 1 step: 7, loss is 1518.2091064453125
epoch: 1 step: 7, loss is 1518.06640625
epoch: 1 step: 8, loss is 1531.715576171875
epoch: 1 step: 9, loss is 1528.3189697265625
epoch: 1 step: 10, loss is 1192.0986328125
epoch: 1 step: 11, loss is 1105.7984619140625
epoch: 1 step: 12, loss is 1483.54833984375
epoch: 1 step: 13, loss is 1470.61474609375
epoch: 1 step: 14, loss is 1085.53588671875
epoch: 1 step: 15, loss is 1480.3011474609375
epoch: 1 step: 16, loss is 1391.0362548828125
epoch: 1 step: 17, loss is 1335.4244384765625
```

So you can observe the training log all the time.

3.2.2. Training Results and Evaluation

Model evaluation, enter the following command for evaluation.

```
bash scripts/run_eval_cpu.sh [PATH_CHECKPOINT]
# [PATH_CHECKPOINT] The Model checkpoint file
```

View the evaluation log, the eval.log in the current directory. Enter the command as follows:

```
tail -f eval.log
```

On my computer, the evaluation result is shown in the figure below:


```
root@ecs-6715:/home/work/models/official/audio/DeepSpeech2# tail -f eval.log
Ref: whatever lord chelford said miss brandon received it very graciously and even with a momentary smile
Hyp:
WER: 1.0 CER: 1.0

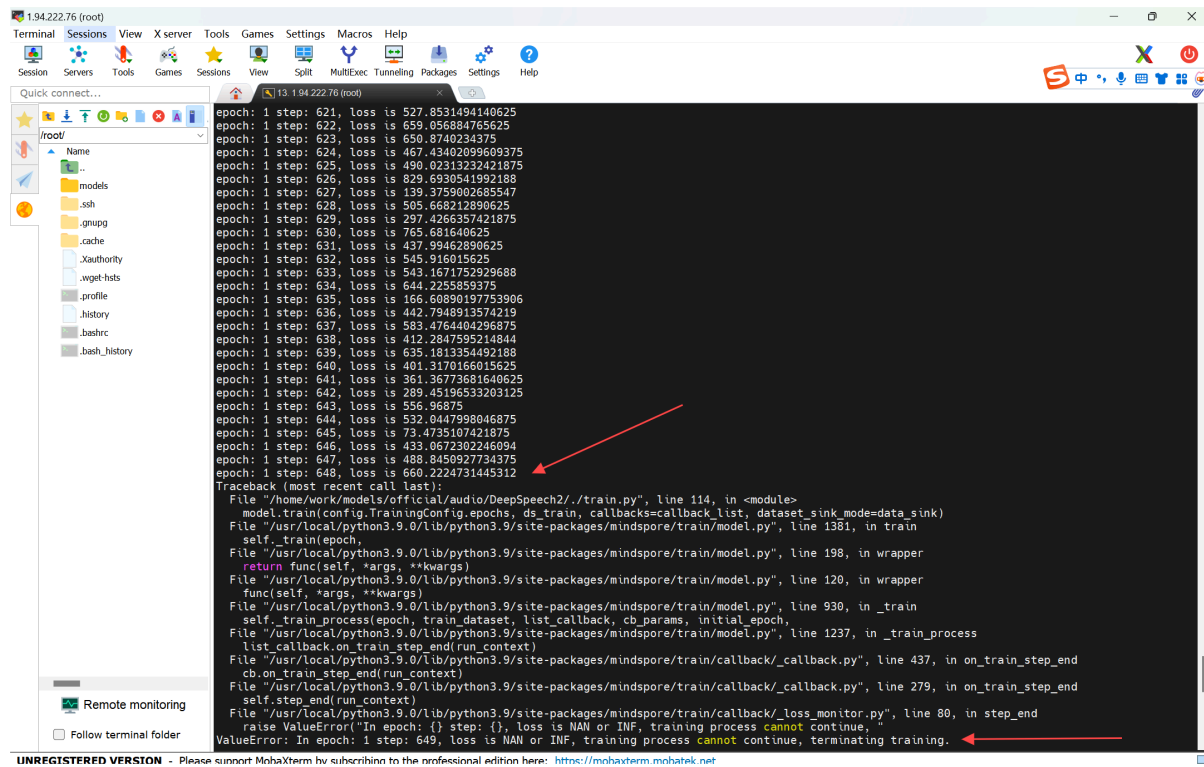
Ref: yes then something better something still grander will surely follow or wherefore should they thus ornament me
Hyp:
WER: 1.0 CER: 1.0
```

ASR refers to the automatic speech recognition technology (Automatic Speech Recognition), which is a technology to convert human speech into text. WER is the word error rate, Word Error Rate (WER) is an important indicator used to evaluate ASR performance, used to evaluate the error rate between predicted text and standard text, so the biggest characteristic of the word error rate is that the smaller, the better.

The two indicators on my computer are probably the same as those above in the experimental manual, and the model evaluation effect is good.

Here I have a problem, when the loss becomes nan or inf after the model training for a long time, it will automatically stop. Just like the following two pictures, I trained twice, one until the step length was more than 500, and the other was more than 600, about half an hour.





```
epoch: 1 step: 621, loss is 527.8531491440625
epoch: 1 step: 622, loss is 659.056884765625
epoch: 1 step: 623, loss is 650.8740234375
epoch: 1 step: 624, loss is 467.43402099609375
epoch: 1 step: 625, loss is 490.02313232421875
epoch: 1 step: 626, loss is 829.6930541992188
epoch: 1 step: 627, loss is 139.3759002685547
epoch: 1 step: 628, loss is 505.668212890625
epoch: 1 step: 629, loss is 297.4266357421875
epoch: 1 step: 630, loss is 765.681640625
epoch: 1 step: 631, loss is 437.99462890625
epoch: 1 step: 632, loss is 545.916015625
epoch: 1 step: 633, loss is 543.1671752929688
epoch: 1 step: 634, loss is 644.2255859375
epoch: 1 step: 635, loss is 166.60890197753906
epoch: 1 step: 636, loss is 442.7948913574219
epoch: 1 step: 637, loss is 583.4764404296875
epoch: 1 step: 638, loss is 412.2847595214844
epoch: 1 step: 639, loss is 635.1813354492188
epoch: 1 step: 640, loss is 401.3170166015625
epoch: 1 step: 641, loss is 361.36773681640625
epoch: 1 step: 642, loss is 289.45196533203125
epoch: 1 step: 643, loss is 556.96875
epoch: 1 step: 644, loss is 532.0447998046875
epoch: 1 step: 645, loss is 73.4735107421875
epoch: 1 step: 646, loss is 433.0672302246894
epoch: 1 step: 647, loss is 488.8450927734375
epoch: 1 step: 648, loss is 660.2224731445312
Traceback (most recent call last):
  File "/home/work/models/official/audio/DeepSpeech2/./train.py", line 114, in <module>
    model.train(config.TrainingConfig.epochs, ds_train, callbacks=callback_list, dataset_sink_mode=data_sink)
  File "/usr/local/python3.9.0/lib/python3.9/site-packages/mindspore/train/model.py", line 1361, in train
    self._train(epoch,
  File "/usr/local/python3.9.0/lib/python3.9/site-packages/mindspore/train/model.py", line 198, in wrapper
    return func(self, *args, **kwargs)
  File "/usr/local/python3.9.0/lib/python3.9/site-packages/mindspore/train/model.py", line 120, in wrapper
    func(self, *args, **kwargs)
  File "/usr/local/python3.9.0/lib/python3.9/site-packages/mindspore/train/model.py", line 930, in _train
    self._train_process(epoch, train_dataset, list_callback, cb_params, initial_epoch,
  File "/usr/local/python3.9.0/lib/python3.9/site-packages/mindspore/train/model.py", line 1237, in _train_process
    list_callback.on_train_step_end(run_context)
  File "/usr/local/python3.9.0/lib/python3.9/site-packages/mindspore/train/callback/_callback.py", line 437, in on_train_step_end
    cb.on_train_step_end(run_context)
  File "/usr/local/python3.9.0/lib/python3.9/site-packages/mindspore/train/callback/_callback.py", line 279, in on_train_step_end
    self.step_end(run_context)
  File "/usr/local/python3.9.0/lib/python3.9/site-packages/mindspore/train/callback/_loss_monitor.py", line 80, in step_end
    raise ValueError("In epoch: {} step: {}, loss is NAN or INF, training process cannot continue, terminating training.")
ValueError: In epoch: 1 step: 649, loss is NAN or INF, training process cannot continue, terminating training.
```

After many times of training, it is inevitable that this situation, I combine the data here to think there are the following several possibilities:

1. The learning rate is too high, which may lead to the gradient explosion, and the parameter update amplitude is too large, which makes the model weight becomes unstable. In the early stage of training, the update range of model weights is large. If the learning rate is too high, it is easy to lead to too large gradient and cause numerical overflow.
2. Gradient explosion, in deep networks the gradient may grow exponentially with backpropagation, leading in numerical spillover. When dealing with long sequences, the gradients can easily accumulate and eventually explode.
3. If the value of the loss function is unstable, the loss function may return a large value or directly return an invalid value. If the output of the model is logits and the value is too large, the numerical overflow may be caused when calculating the Softmax. Cross-entropy loss is prone to extreme values with probability approaching 0 or 1.
4. Sequence length or batch size problems, increasing memory requirements and computational complexity when handling long speech sequences or large volumes of data, may trigger numerical overflows. The longer the sequence length, the more information needed to process, easily leading to gradient problems. When the batch size is too large, the gradient of each update may be too intense.

3.3. Model Export

Model export that requires the following code to modify the export.py file.

```
config = train_config
context.set_context(mode=context.GRAPH_MODE, device_target="CPU", save_graphs=False)
with open(config.DataConfig.labels_path) as label_file:
    labels = json.load(label_file)

config = train_config
context.set_context(mode=context.GRAPH_MODE, device_target="CPU", save_graphs=False)
with open(config.DataConfig.labels_path) as label_file:
    labels = json.load(label_file)
```

Enter the following command for converting and exporting the model file.

```
root@ecs-6715:/home/work/models/official/audio/DeepSpeech2# python export.py --pre_trained_model_path /home/work/models/official/audio/DeepSpeech2/DeepSpeech.ckpt
Successfully loading the pre-trained model
```

Then enter the command `ll` in the directory, and you can see the model file:

```
root@ecs-6715:/home/work/models/official/audio/DeepSpeech2# ll
total 1873172
drwxr-xr-x 6 root root      4096 Dec 14 13:37 ./
drwxr-xr-x 7 root root      4096 Dec 13 14:27 ../
drwx----- 2 root root      4096 Dec 14 13:20 checkpoint/
-r----- 1 root root    346524583 Dec 14 13:45 deepspeech2.mindir
-rw-r--r-- 1 root root 1039306262 Jan 10 2023 DeepSpeech.ckpt
drwxr-xr-x 3 root root      4096 Dec 13 18:16 deepspeech_pytorch/
-rw-r--r-- 1 root root      8216 Dec 14 02:36 eval.log
-rw-r--r-- 1 root root     5690 Dec 13 14:27 eval.py
-rw-r--r-- 1 root root     3003 Dec 14 13:37 export.py
-rw-r--r-- 1 root root       205 Dec 13 14:27 labels.json
-rw-r--r-- 1 root root 532178606 Dec 14 01:12 librispeech_val_output.bin
-rw-r--r-- 1 root root     4765 Dec 13 14:27 quick_start.py
-rw-r--r-- 1 root root    13387 Dec 13 14:27 README-CN.md
-rw-r--r-- 1 root root   14740 Dec 13 14:27 README.md
-rw-r--r-- 1 root root        33 Dec 13 14:27 requirements.txt
drwxr-xr-x 2 root root      4096 Dec 13 14:27 scripts/
drwxr-xr-x 3 root root      4096 Dec 13 18:13 src/
-rw-r--r-- 1 root root     3258 Dec 14 13:20 train.log
-rw-r--r-- 1 root root     5442 Dec 13 14:27 train.py
```

4. Answer the Question

4.1. Question Description

1 round of model training, need 50 hours, how to improve the speed?

4.2. Problem Thinking

- **Replace the stronger hardware**

Using either a GPU or a TPU.

GPU: When training a deep learning model, the GPU is the most commonly used acceleration hardware. Compared with CPU, GPU can greatly improve the training speed when processing matrix operations. A NVIDIA GPU, such as A100, V100, or RTX 3090, is recommended.

TPU: If you use the Google Cloud, or any other TPU-enabled platform, the TPU performance may be stronger and is suitable for large-scale deep learning tasks.

Multi-GPU parallel training: If a single GPU is not enough, multiple GPU can be used for parallel training (data parallel or model parallel). Frameworks such as TensorFlow, PyTorch, and Mindspore all support this training approach.

Using a high-performance server:

If train on a cloud server, consider using stronger instances (such as NVIDIA Tesla V100, A100, or TPU instance). Cloud computing platforms: AWS, Google Cloud, and Azure all provide powerful computing resources that can be expanded as needed.

- **Training with a mixed-precision approach**

FP16 (Hybrid Precision Training): Hybrid precision training (FP16) can accelerate training while saving memory. It uses most of the calculations using 16-bit floating-points instead of 32-bit, thereby increasing computation speed and reducing memory footprint.

- **Data preprocessing and loading optimization**

Data preprocessing acceleration: If the data preprocessing before the training (such as audio feature extraction, image enhancement, etc.) is too time-consuming, it will significantly affect the training speed.

- **Adjust the bulk size**

Increase the batch size (Batch Size): Increasing the batch size can reduce the time required to calculate each step. However, increasing the batch size increases memory consumption and needs to be adjusted for hardware resources.

Increasing batch size: Sometimes a large batch size can cause video memory outages. In this case, the "progressive batch size" strategy can be adopted to gradually increase the batch size until the video memory reaches the maximum capacity.

- **Distributed training**

Data parallel: use multiple GPU or computing nodes, using data parallel training. Each compute node or GPU processes a different batch of the training data to accelerate the training by synchronously updating the model parameters.

Hybrid parallelism: combine model parallelism and data parallelism, distribute the model to multiple devices, and accelerate the training through data parallelism. This applies to very large models.

- **Model architecture optimization**

Pruning (Pruning): Pruning is a technique that reduces the number of model parameters to accelerate reasoning and training by removing redundant neurons or connections.

Knowledge distillation (Knowledge Distillation): Train a small model to simulate the output of a large pre-training model to accelerate the inference and training process.

Network compression: smaller networks (such as MobileNet, EfficientNet) are used to replace complex network models. For some tasks, using a lighter architecture significantly reduces the training time.